

# AWS multi-tenant

▼ Upgrade EKS

## What This Script Does [↗](#)

This script automates the installation and upgrade of **Karpenter** on an **Amazon EKS (Elastic Kubernetes Service)** cluster using **Helm**.

## What is Karpenter? [↗](#)

Karpenter is an **autoscaler** for Kubernetes that automatically provisions and deprovisions EC2 instances based on workload demands. Unlike the traditional **Cluster Autoscaler**, Karpenter directly interacts with EC2, making scaling decisions faster and more efficient.

## Script Breakdown [↗](#)

The script:

1. **Finds an existing EKS cluster**
2. **Finds necessary AWS IAM roles and resources**
3. **Deploys/Upgrades Karpenter using Helm**
4. **Configures Karpenter with proper IAM permissions and Kubernetes settings**
5. **Deploys additional Karpenter components (EC2NodeClass, NodePool)**

## Step-by-Step Execution [↗](#)

### 1. Load Environment Variables & Functions [↗](#)

```
1 . ${WORKSPACE}/scripts/functions.sh
2
```

This loads helper functions (like `plog` for logging and `die` for error handling) from a script.

### 2. Enable Debug Mode (Optional) [↗](#)

```
1 debug=false
2 if [ "$1" == "debug" ];then
3     debug=true
4 fi
5
```

If the script is run with `debug` mode ( `./script.sh debug` ), it enables Helm's `--debug` and `--dry-run` flags.

### 3. Get the EKS Cluster Name [↗](#)

```
1 cluster_name=$(aws eks list-clusters | jq -r '.clusters[0]')
2 if [ -z "${cluster_name}" ];then
3     die "Can't get cluster_name"
4 fi
5
```

- Retrieves the **first** EKS cluster in AWS using `aws eks list-clusters`.
- If no cluster is found, it **exits with an error**.

✅ **Why?** Karpenter **must** be installed in an existing EKS cluster.

### 4. Get the Karpenter IAM Role [↗](#)

```
1 karpenter_role=$(aws iam list-roles | jq -r '[][].Arn' | grep KarpenterController)
2 if [ -z "${karpenter_role}" ];then
3     die "Can't get karpenter_role"
4 fi
5
```

- Searches for an **IAM Role** with `KarpenterController` in its name.
- If no role is found, the script **exits**.

✅ **Why?** Karpenter requires an IAM Role to interact with EC2 and autoscale nodes.

### 5. Get the Karpenter Node IAM Role [↗](#)

```
1 karpenter_role_name=$(aws iam list-roles | jq -r '[][].RoleName' | grep Karpenter-eks)
2 if [ -z "${karpenter_role_name}" ];then
3     die "Can't get karpenter_role_name"
4 fi
5
```

```
4 fi
5
```

- Finds the IAM Role Name for worker nodes managed by Karpenter.

✔ **Why?** Each Karpenter-provisioned node needs an **IAM Role** with EC2 permissions.

## 6. Set Default Versions for Karpenter [↗](#)

```
1 if [ -z "${version_tag}" ];then
2     version_tag="0.35.4"
3 fi
4
5 if [ -z "${version_digest}" ];then
6     version_digest="sha256:d59b5d4c58011615b593dd80f115106efad034104c2d7f016c01c19d8e0b21d9"
7 fi
8
9 if [ -z "${image_repo}" ];then
10     image_repo="690359577697.dkr.ecr.us-east-1.amazonaws.com/infra"
11 fi
12
13 if [ -z "${chart_version}" ];then
14     chart_version="0.36.0"
15 fi
16
```

- Defines **default versions** for:
  - Karpenter Helm Chart
  - Karpenter Docker Image
  - Image repository location

✔ **Why?** Allows easy upgrades by modifying **chart\_version** and **image\_tag**.

## 7. Get the Karpenter SQS Queue [↗](#)

```
1 karpenter_queue=$(aws sqs list-queues | jq -r '.QueueUrls[]' | grep Karpenter | awk -F\ ' {print $NF}')
2 if [ -z "${karpenter_queue}" ];then
3     die "Can't get karpenter_queue"
4 fi
5
```

- Retrieves the **SQS queue name** for Karpenter.
- This queue handles **interruption events** (e.g., when a spot instance is about to be terminated).

✔ **Why?** Karpenter needs an **SQS queue** to gracefully handle instance terminations.

## 8. Get the Cluster Endpoint [↗](#)

```
1 cluster_endpoint=$(aws eks describe-cluster --name ${cluster_name} | jq -r '.cluster.endpoint')
2 if [ -z "${cluster_endpoint}" ];then
3     die "Can't get cluster_endpoint"
4 fi
5
```

- Retrieves the **API endpoint** of the EKS cluster.

✔ **Why?** Karpenter must know the **Kubernetes API** location to manage nodes.

## 9. Create a Service-Linked Role for Spot Instances [↗](#)

```
1 aws iam create-service-linked-role --aws-service-name spot.amazonaws.com > /dev/null 2>&1
2
```

- Ensures AWS has a **service-linked role** for managing **Spot instances**.

✔ **Why?** Karpenter can provision **Spot Instances** to save costs.

## 10. Install/Upgrade Karpenter Using Helm [↗](#)

```
1 helm upgrade --install ${debug_info} karpenter oci://public.ecr.aws/karpenter/karpenter \
2     --version "${chart_version}" --namespace "karpenter" --create-namespace \
3     --set "serviceAccount.annotations.eks.amazonaws.com/role-arn=${karpenter_role}" \
4     --set "logLevel=debug" \
5     --set "settings.clusterName=${cluster_name}" \
6     --set "settings.clusterEndpoint=${cluster_endpoint}" \
7     --set "settings.interruptionQueue=${karpenter_queue}" \
8     --set "controller.image.repository=${image_repo}" \
```

```

9  --set "controller.image.tag=karpenter-controller-${version_tag}" \
10 --set "controller.image.digest=${version_digest}" \
11 --set controller.resources.requests.cpu=1 \
12 --set controller.resources.requests.memory=1Gi \
13 --set controller.resources.limits.cpu=1 \
14 --set controller.resources.limits.memory=1Gi \
15 --wait
16

```

- Installs (or upgrades) Karpenter using **Helm**.
- Configures it with:
  - The **IAM Role**
  - The **Cluster API endpoint**
  - The **SQS queue**
  - The **container image**
  - CPU & memory limits

✅ **Why?** Helm simplifies managing Karpenter.

## 11. Install/Upgrade EC2NodeClass and NodePool [↗](#)

```

1 helm upgrade --install karpenter-ec2nodeclass karpenter_charts/EC2NodeClass --set "karpenter_node_iam_role_name=${karpenter_role_name}"
2 helm upgrade --install karpenter-nodepool karpenter_charts/NodePool
3

```

- **EC2NodeClass** defines what type of EC2 instances Karpenter can provision.
- **NodePool** defines scaling rules.

✅ **Why?** These are **essential components** for managing EC2 nodes in Kubernetes.

## 12. Verify the Installation [↗](#)

```

1 helm list -A --time-format "2006-01-02 15:04"
2

```

- Lists installed Helm charts.

✅ **Why?** Confirms whether **Karpenter was successfully installed**.

## What Does This Script Achieve? [↗](#)

- ✅ **Automates the installation & upgrade** of Karpenter.
- ✅ **Ensures AWS dependencies (IAM, SQS, EKS) are set up**.
- ✅ **Deploys Karpenter with proper configurations for Kubernetes autoscaling**.
- ✅ **Allows easy version updates** using Helm.

## Explanation of Your EC2NodeClass Configuration [↗](#)

This **EC2NodeClass** resource defines how Karpenter provisions **EC2 instances** in your Amazon EKS cluster. It controls **AMI selection, storage, networking, security groups, and user data configuration**.

## Step-by-Step Breakdown [↗](#)

### 1. Resource Type and API Version [↗](#)

```

1 apiVersion: karpenter.k8s.aws/v1beta1
2 kind: EC2NodeClass
3 metadata:
4   name: default
5

```

- **apiVersion:** `karpenter.k8s.aws/v1beta1` → This is a Karpenter-specific **CRD (Custom Resource Definition)**.
- **kind:** `EC2NodeClass` → Defines how Karpenter provisions new EC2 instances.
- **metadata.name:** `default` → This NodeClass is named **default**, meaning it may be referenced by NodePools.

✅ **Why?** This **NodeClass** acts as a blueprint for provisioning nodes.

### 2. Specify the AMI (Amazon Machine Image) [↗](#)

```

1 spec:
2   amiFamily: AL2
3

```

- **AL2** refers to **Amazon Linux 2**.

- Karenter automatically selects an Amazon-provided **optimized AMI**.

✔ **Why?** Ensures nodes are using a stable and optimized **Kubernetes-compatible** OS.

### 3. Custom User Data (Startup Script) [↗](#)

```
1  userData: |
2    #!/bin/bash
3    fallocate -l 32G /swapfile
4    chmod 600 /swapfile
5    mkswap /swapfile
6    swapon /swapfile
7    sysctl -w vm.swappiness=10
8    jq '.failSwapOn=false' /etc/kubernetes/kubelet/kubelet-config.json | jq '.featureGates.NodeSwap=true' | jq '.memorySwap.swapBehavior="UnlimitedSwap"' >
   /etc/kubernetes/kubelet/kubelet-config.json.new
9    cp -rf /etc/kubernetes/kubelet/kubelet-config.json /etc/kubernetes/kubelet/kubelet-config.json.org
10   cp -rf /etc/kubernetes/kubelet/kubelet-config.json.new /etc/kubernetes/kubelet/kubelet-config.json
11
```

This **userData** script runs when the instance starts:

#### 1. Creates a 32GB Swap File

```
1  fallocate -l 32G /swapfile
2  chmod 600 /swapfile
3  mkswap /swapfile
4  swapon /swapfile
5  sysctl -w vm.swappiness=10
6
```

- **Creates a 32GB swap file** to expand virtual memory.
- **Sets** `swappiness=10` to reduce aggressive swap usage.

#### 2. Enables Kubernetes NodeSwap

```
1  jq '.failSwapOn=false' /etc/kubernetes/kubelet/kubelet-config.json | jq '.featureGates.NodeSwap=true' | jq '.memorySwap.swapBehavior="UnlimitedSwap"' >
   /etc/kubernetes/kubelet/kubelet-config.json.new
2
```

- Modifies Kubelet's configuration to **allow swap memory**.
- Enables the `NodeSwap` feature gate.
- Sets **unlimited swap behavior**.

✔ **Why?** This ensures nodes **do not fail due to memory pressure** and allows swap usage.

### 4. IAM Role for Node Instances [↗](#)

```
1  role: {{ .Values.karpenter_node_iam_role_name }}
2
```

- Uses the IAM role defined in Helm ({{ .Values.karpenter\_node\_iam\_role\_name }}).
- The role allows EC2 instances to join the EKS cluster and interact with AWS services.

✔ **Why?** Ensures that nodes **inherit IAM permissions** needed for Kubernetes operations.

### 5. Define EC2 Instance Storage [↗](#)

```
1  blockDeviceMappings:
2    - deviceName: /dev/xvda
3      ebs:
4        volumeSize: {{ .Values.ec2_node_class_volume_size }}
5        volumeType: {{ .Values.ec2_node_class_volume_type }}
6        iops: {{ .Values.ec2_node_class_volume_iops }}
7        encrypted: true
8        deleteOnTermination: true
9        throughput: {{ .Values.ec2_node_class_throughput }}
10
```

- **Device Name** `/dev/xvda` → This is the primary volume.
- **Volume Size** → {{ .Values.ec2\_node\_class\_volume\_size }} (Defined in Helm).
- **Volume Type** → {{ .Values.ec2\_node\_class\_volume\_type }} (e.g., `gp3` or `io1`).
- **IOPS (Input/Output Operations Per Second)** → {{ .Values.ec2\_node\_class\_volume\_iops }}.
- **Throughput** → {{ .Values.ec2\_node\_class\_throughput }} (Important for `gp3` volumes).
- **Encrypted:** `true` → Ensures data security.
- **Delete on Termination:** `true` → Deletes the volume when the instance is terminated.

✔ **Why?** Controls **storage configuration**, ensuring performance and security.

6. Select AWS Subnets for Node Deployment

```
1 subnetSelectorTerms:
2   - tags:
3     karpenter.sh/discovery: eks
4
```

- Selects **subnets** with the tag `karpenter.sh/discovery: eks`.
- ✔ **Why?** Ensures instances **launch in the correct VPC subnets**.

7. Assign Security Groups to Nodes

```
1 securityGroupSelectorTerms:
2   - tags:
3     karpenter.sh/discovery: eks
4
```

- Selects **security groups** based on the same `karpenter.sh/discovery: eks` tag.
- ✔ **Why?** Ensures worker nodes **inherit the correct security settings**.

8. Apply Tags to Instances

```
1 tags:
2   karpenter.sh/discovery: eks
3
```

- Tags all created instances with `karpenter.sh/discovery: eks`.
- ✔ **Why?** Helps Karpenter **track and manage instances**.

Summary: What Does This EC2NodeClass Do?

Feature	Description
AMI Selection	Uses Amazon Linux 2 (AL2).
Startup Script	Creates a <b>32GB swap file</b> and enables <b>NodeSwap</b> .
IAM Role	Uses a dynamically assigned IAM role for worker nodes.
EBS Storage	Defines <b>storage size, type, IOPS, throughput, and encryption</b> .
Subnet Selection	Chooses subnets with the tag <code>karpenter.sh/discovery: eks</code> .
Security Groups	Assigns security groups with the tag <code>karpenter.sh/discovery: eks</code> .
Tags	Labels EC2 instances for tracking.

Why is This EC2NodeClass Important?

- ♦ Controls how Karpenter provisions EC2 nodes.
- ♦ Ensures nodes have the right IAM, networking, and storage settings.
- ♦ Improves performance with swap memory and storage optimizations.
- ♦ Allows Kubernetes workloads to scale efficiently.

Explanation of Your Karpenter NodePool Configuration

This **NodePool** resource defines **how Karpenter provisions and manages EC2 instances** in your Kubernetes cluster. It controls **which instance types are allowed, where nodes can be created, autoscaling limits, and node lifecycle policies**.

Step-by-Step Breakdown

1. Define Resource Type and Metadata

```
1 apiVersion: karpenter.sh/v1beta1
2 kind: NodePool
3 metadata:
4   name: default
5
```

- **apiVersion:** `karpenter.sh/v1beta1` → This is a Karpenter-specific **CRD (Custom Resource Definition)**.
- **kind:** `NodePool` → Defines a **scaling policy** for Kubernetes nodes.
- **metadata.name:** `default` → The name of this NodePool is **default**.

✓ **Why?** This **NodePool** tells Karpenter how to provision EC2 nodes.

## 2. Define Node Template (How Nodes Are Created) [↗](#)

```
1 spec:
2   template:
3     metadata:
4       labels:
5         created_by: karpenter
6
```

- **Labels** → Assigns the label `created_by: karpenter` to each node.

✓ **Why?** This makes it **easy to filter** and manage Karpenter-created nodes.

## 3. Associate with an EC2NodeClass [↗](#)

```
1   spec:
2     nodeClassRef:
3       name: default
4
```

- **nodeClassRef:** `default` → Links this **NodePool** to an **EC2NodeClass** named `default`.

✓ **Why?** The **EC2NodeClass** defines AMI, storage, IAM roles, and network settings.

## 4. Define Node Scheduling Constraints [↗](#)

```
1   requirements:
2     - key: "node.kubernetes.io/instance-type"
3       operator: In
4       values: {{ .Values.karpenter_instance_types | toJson }}
5
```

- **Instance Type Constraint**
  - This restricts Karpenter to **specific EC2 instance types**.
  - `{{ .Values.karpenter_instance_types | toJson }}` dynamically pulls values from Helm.

Example (after Helm resolves values):

```
1   requirements:
2     - key: "node.kubernetes.io/instance-type"
3       operator: In
4       values: ["t3.medium", "m5.large", "c5.xlarge"]
5
```

✓ **Why?** Ensures that **only approved instance types** are used.

```
1     - key: "topology.kubernetes.io/zone"
2       operator: In
3       values: {{ .Values.azs | toJson }}
4
```

- **Availability Zone Constraint**
  - Restricts nodes to **specific AWS Availability Zones**.
  - Example values after Helm resolves:

```
1   values: ["us-east-1a", "us-east-1b", "us-east-1c"]
2
```

✓ **Why?** Ensures workloads **only run in specified zones**.

```
1     - key: "karpenter.sh/capacity-type"
2       operator: In
3       values: {{ .Values.karpenter_instance_capacity_types | toJson }}
4
```

- **Capacity Type Constraint**
  - Controls whether Karpenter uses:
    - **"on-demand"** instances.
    - **"spot"** instances (cheaper but can be terminated anytime).
  - Example values after Helm resolves:

```
1   values: ["on-demand", "spot"]
```

✓ **Why?** Allows cost-optimization by mixing **Spot and On-Demand instances**.

## 5. Set CPU Limits for Auto-Scaling [↗](#)

```
1 limits:
2   cpu: 1000
3
```

- **Limits total vCPUs across all nodes** managed by this **NodePool**.
- **1000 CPU cores** → Limits the maximum cluster size.

✓ **Why?** Prevents **uncontrolled auto-scaling** from over-provisioning nodes.

## 6. Define Node Disruption and Consolidation Policies [↗](#)

```
1 disruption:
2   consolidationPolicy: WhenEmpty
3   consolidateAfter: 30s
4   expireAfter: 'Never'
5
```

- `consolidationPolicy: WhenEmpty` →
  - Karpenter **removes underutilized nodes only when they are empty**.
- `consolidateAfter: 30s` →
  - Karpenter checks for **unused nodes every 30 seconds**.
- `expireAfter: 'Never'` →
  - Nodes never expire unless scaled down manually or consolidated.

✓ **Why?** Prevents **unnecessary node deletions** while still optimizing resource usage.

## Summary: What Does This NodePool Do? [↗](#)

Feature	Description
<b>Instance Type Selection</b>	Restricts instances to approved types.
<b>Availability Zones</b>	Ensures nodes are provisioned only in selected AZs.
<b>Capacity Type</b>	Supports <b>Spot</b> and <b>On-Demand</b> instances.
<b>Max CPU Limit</b>	Prevents excessive node scaling beyond <b>1000 CPUs</b> .
<b>Consolidation</b>	Automatically removes <b>empty</b> nodes after <b>30 seconds</b> .
<b>Expiration</b>	Nodes <b>never expire</b> automatically.

## Why is This NodePool Important? [↗](#)

- **Controls auto-scaling behavior** for Kubernetes worker nodes.
- **Optimizes costs** by allowing a mix of **Spot** and **On-Demand** instances.
- **Improves scheduling** by defining **which instance types and AZs** are allowed.
- **Reduces waste** by consolidating **idle nodes** after 30s.

### ✓ ECR lifecycle policies for unused images

AIM: remove unused images from the ECR

- **PRODUCTION\_IMAGE\_TTL = 90** (Giddy\_... , Felix\_...)
- **FB\_IMAGE\_TTL = 60** ( \*\_CR\_FB\_\* )
- **MASTER\_IMAGE\_TTL = 60** (master) - step 2.1
- **PRODUCTION\_IMAGE\_WITH\_GA\_BUILD** -> **Not deleting GA images** (Were released as GA) - step 2.4
- **UNKOWN\_IMAGE\_PATTERN\_TTL = 90** (??, need to take a look at the code)

Step 1: create a list of all current image in the ECR

Skybox's ECR is residing in "skybox-shared-services" account.

There are a few jobs that update the images that are in the ECR:

- Build\_Saas\_Agents - [https://jenkins-srv/job/Build\\_Saas\\_Agents/](https://jenkins-srv/job/Build_Saas_Agents/) - MySQL, Elastic, Artemis, Cloudwatch, Datadog
- non\_production\_build - [https://jenkins-srv/job/non\\_production\\_build/](https://jenkins-srv/job/non_production_build/)

- Version's pipeline - <https://jenkins-srv/job/Giddy/> - Giddy, felix..

I created a script in devopstols/aws\_scripts/shared\_account\_utils/get\_all\_ecr\_images.sh that returns a list of all the images that are in all the repositories within the ECR.

#### Step 2: handle all "do-not-delete" images

##### Step 2.1: life cycle policies

Each repository within the ECR has its own life cycle policy (see end of section for the policies as of 11-feb-2025).

The lifecycle policies handle images with the prefixes - [SKY-] , [master\_] , [cherry-pick-] [collector-base] , [elastic-SKY-] , [mysql-SKY-] , [artemis-SKY-] , [cloudwatch-SKY-] , [datadog-SKY-] , [elastic-master\_] , [mysql-master\_] , [artemis-master\_] , [cloudwatch-master\_] , [datadog-master\_]

##### Step 2.2: create a "white list" of all currently used by an ECS images

We need to handle the images that are currently in production. To do this we need to create a "white list" of all images that are currently being used by an ECS of an account.

To see image used by ECS = AWS → ECS → task definition

I created a script in devopstols/aws\_scripts/shared\_account\_utils/get\_images\_used\_by\_ecs.sh that returns a list of all images that are being used by an ECS (this needs to run for each account).

##### Step 2.3: create a "white list" of all currently used by an EKS images

f

##### Step 2.4: get the list of GA versions

GA images are not to be deleted!

##### Step 3: create a list of images to delete

In step 2.1 we got a list of prefixes that we can ignore, as the life cycle policy handles them. In step 2.2 we got a white list of images that are not to be deleted as they are currently in use. and so, we need to subtract the images from step 2 from the list of all images within the ECR that we got in step 1.

#### **Repositories with No Lifecycle Policy**

- development
- ecr-public/eks-distro/kubernetes/pause
- skyboxdev/app
- skyboxdev/services/service-check-access
- skyboxreleases/app
- skyboxreleases/collector
- skyboxreleases/server

---

#### **Repositories with Lifecycle Policies**

##### **Short-lived images (1-day expiry)**

- skyboxdev/services/service-vtm-service
- skyboxdev/services/service-model-service
- skyboxdev/services/service-hello-world
- skyboxdev/services/service-conga
  - **Rule 1:** Expire **all images** after **1 day**

---

##### **SkyboxDev Server**

- **Rule 1:** Expire **untagged images** after **1 day**
- **Rule 2:** Expire **tagged images** with prefix [SKY-] after **14 days**
- **Rule 3:** Expire **tagged images** with prefix [master\_] after **14 days**
- **Rule 4:** Expire **tagged images** with prefix [cherry-pick-] after **1 day**

---

##### **SkyboxDev Collector**

- **Rule 1:** Expire **untagged images** after **1 day**
- **Rule 2:** Expire **tagged images** with prefix [SKY-] after **7 days**
- **Rule 3:** Expire **tagged images** with prefix [master\_] after **14 days**
- **Rule 4:** Expire **tagged images** with prefix [cherry-pick-] after **1 day**

---

##### **Infra**

- **Rule 1:** Expire **untagged images** after **1 day**
- **Rule 2:** Expire **tagged images** with prefix [collector-base] after **1 day**

---

##### **DevOps Agent**

- **Rule 1:** Expire **untagged images** after **1 day**



- **Rules 2-6:** Expire **tagged images** with prefix `[elastic-SKY-]`, `[mysql-SKY-]`, `[artemis-SKY-]`, `[cloudwatch-SKY-]`, `[datadog-SKY-]` after **14 days**
- **Rules 7-11:** Expire **tagged images** with prefix `[elastic-master_]`, `[mysql-master_]`, `[artemis-master_]`, `[cloudwatch-master_]`, `[datadog-master_]` after **14 days**

#### ▼ Deploying with Terraform across an account

- change the necessary files in `mt-deployment/`. for example:
  - `variables.tf` - insert within the map `"bash_scripts_to_s3"` an additional bash script
  - `ssm.tf` - create a new SSM document
- Run the pipeline named **"aws\_eks\_account\_provisioning"** **for each wanted account** (no way to perform this action on multiple accounts for safety reasons)
  - The pipeline will wait for you to approve the `plan.json` (can be found in the archived artifacts of the specific Jenkins job). Visualize the file through <https://sb-tfui> and check:
    - If `"cluster_creation"` is showing go to AWS access portal → account → global/admin credentials → EKS → access → role-global-terraform → delete
    - If anything related to network shows know you will have to re-run the pipeline until they won't show (the exception is - `module.tgw_attachment.aws_ec2_transit_gateway_vpc_attachment.local_account`)
  - Re-run the pipeline until you don't see updates/creations that are not wanted and/or no resources added.

#### ▼ Delete PVC in multi-tenant AWS EKS - elasticsearch + mysql + data

AIM: delete resources connected to a stateful set (mysql/elasticsearch/data)

Each tenant has its own stateful set, with replica set to 1.

connection to AWS -activated by `"session_to_eks.sh"` (creates an SSH tunnel through a bastion EC2 instance, but using an AWS document called `'port_forward_to_remote_document'`)

Pipeline - `aws_mt_wipedb`

- In Jenkins
  - create a new item by copying `"aws_eks_tenant_resize"` [📺 How To Clone a Jenkins Job](#)
  - remove unnecessary parameters by clicking the `"configure"` button and scrolling down until you reach an unwanted button.
  - if you need to add new parameters [📺 How Do I Add a Choice Parameter in Jenkins?](#)
- The Jenkinsfile
  - based on `aws_eks_tenant_resize/Jenkinsfile`

Script - `aws_mt_wipedb.sh`

- stop stateful set by lowering replica to zero
- delete stateful set
- delete pvc
  - `mysql = data-data-elastic-skybox-0`
  - `elasticsearch = data-mysql-skybox-0`
  - `data = skybox-data-skybox-0`
- delete files from the bastion of the customer's cluster. In path `"cd /efs/customers/<tenant name>"` search the latest `"confx"` directory and delete the following files:
  - for `wipe_db`

```
1 rm -f \"$latest_confx/*.pem\"
2 rm -f \"$latest_confx/*.cksum\"
3 rm -f \"$latest_confx/*.p12\"
```
  - for elastic search
 

```
1 rm -f \"$latest_confx/server.skybox.elastic.crt\"
2 rm -f \"$latest_confx/server.skybox.elastic.key\"
```
- recreate stateful - by calling pipeline `"aws_eks_tenant_provisioning_update"` )
  - search for usage examples in devopstool, if not found - [📺 Getting started with Pipeline](#)
  - Snippet Generator - steps for a Scripted Pipeline | `'steps'` block in a `'stage'` in a Declarative Pipeline
  - Declarative Directive Generator - sections and directives used to define a Declarative Pipeline.

Test the pipeline

- [📺 AWS access portal](#) - access with `"skybox-internal-devops03s-dev"` to `"skybox-internal-devops03s-dev"`
  - make sure you are at one of the following:
    - N. virginia - us-east-1
    - singapore - ap-southeast-1
    - frankfurt - eu-central-1
  - EC2 - micro is the bastion that `delete_confx_files.sh` accesses to delete file -within the instance
    - access k9 = `/efs/k9s`
    - healthcheck = `bash -x /healthcheck.sh`
- [https://sb-consul/ui/dc1/kv/aws/ou/non-production-rnd/jenkins\\_multi\\_tenant/003857080240/devopscust01/skybox\\_public\\_url/edit](https://sb-consul/ui/dc1/kv/aws/ou/non-production-rnd/jenkins_multi_tenant/003857080240/devopscust01/skybox_public_url/edit) - the url for devopscust01 within the multi tenant of devops3
  - ensure you can login to it

- username = saasadmin
- password = [AWS Systems Manager](#) → Parameter Store → /devopscust01/production\_saasadmin\_password