

# Grafana

✓ Create Grafana Dashboard to Monitor Skybox/TF Version and Tenant Status

## Objective:

- Build a pipeline to collect and store information about tenants in a multi-tenant AWS environment.
- Visualize the collected data in a Grafana dashboard.

## Components:

### 1. Pipeline - aws\_mt\_inventory

- a. Get mapping of account\_id-account\_name from AWS - - using `aws organizations list-accounts | jq -r '.Accounts[]'` with root access.
- b. Get a list of all multi-tenant from consul
- c. Sub-pipeline - aws\_mt\_account\_inventory - runs in parallel for all accounts
  - i. Connect to AWS to get the stateful-sets of the tenants of the account
  - ii. Script (aws\_mt\_inventory.py) - collects all the data needed and send to elastic search with a unique id per each tenant.
- d. Pipeline - aws\_eks\_tenant\_provisioning - already exists, just add an upsert to a document, to ensure consistent data from multiple sources.
  - i. if tenant is destroyed - change its EKS status to deleted.

### 2. Elasticsearch Document:

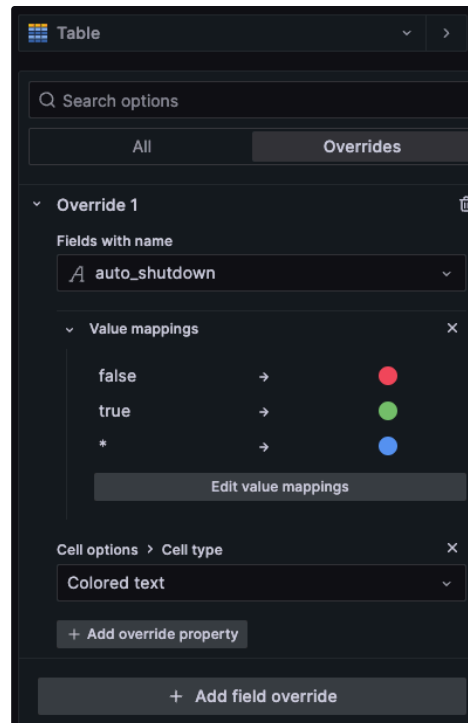
#### ◦ Fields:

- customer (string): Name of the customer.
- ou (string): Organizational Unit.
- account\_id (string): AWS account ID.
- aws\_account\_name (string): Name of the AWS account.
- terraform\_version (string): Terraform version used for provisioning.
- connected\_to\_voltra (boolean): Indicates connection to Voltra.
- skybox\_version (string): Skybox version (tenant and consul tag).
- is\_url\_reachable (boolean): Indicates if the URL is reachable.
- stateful\_set\_status (string): Status of the stateful set.
- auto\_shutdown (boolean)
- ebs\_size (string)
- enable\_cloudwatch\_agent (boolean)
- enable\_datadog\_agent (boolean)
- env (string)
- size (string)
- skybox\_public\_url (string)
- tag (string)

### 3. Grafana Dashboard (AWS Multi-Tenant Summary):

- **Data Source:** Elasticsearch indices containing "aws\_mt\_inventory\*" data.
- **Query:** Filters data based on customer, OU, account ID, and account name.
  - query: customer: \$customer AND ou: \$ou AND account\_id: \$account\_id AND aws\_account\_name: \$aws\_account\_name AND \_index:"aws\_mt\_inventory"
  - index:"aws\_mt\_inventory\_ownership"
- **Transformations:**
  - Group by transformation -
    - Groups data by customer (primary grouping).
    - Calculates "first\*" for each field to display the latest non-null value per customer.
    - Ignored for column not to be displayed.
  - Organize fields transformation - organizes fields with desired column names + order columns.
- **Display:**
  - Table format.

- `consul_client.put(f"`  
`{AWS_OU_PREFIX}/{ou}/{AWS_MT_CONFIGURATION_SUFFIX}/{account}/{customer}/{MT_ACCOUNTS_AUTO_SHUTDOWN_SUFFIX}",`  
`custom_control_value)`
- Hides unnecessary columns.
- Applies color coding to specific columns based on value patterns.
  - add override property → value mapping (does not exist for this client, `^(?!true|false).*` )
  - add override property → cell options > cell type



- ✓ Update doc in elasticsearch to increase Grafana dashboard capabilities

AIM: take the current index and add parameters to it to create a more robust Grafana dashboard.

#### Step 1: Search for Usage of `junit_parser` in Code [🔗](#)

1. Open **PyCharm**.
2. Use the shortcut `Command + Shift + F` to search for `junit_parser` across all files.
3. Update all references where documents are pushed to Elasticsearch through it:
  - `pipeline` = source to target + `NightlyRun`
  - Replace **index name** in `parse_results` with:

```
es_api.bulk_insert("junit_yalmaliah_test", docs, elastic_password=elastic_password)
```

#### Step 2: Modify `junit_parser.py` [🔗](#)

1. **Add Arguments:**
  - Add short and long argument names.
  - Follow the function call trail to ensure proper integration.
2. **Update Index:** Ensure the target index name (`junit_yalmaliah_test`) is used in the `parse_results` function.

#### Step 3: Integrate Updates into the Pipeline [🔗](#)

1. **Add Arguments in the Pipeline:**
  - Through the script (`sh`): Add arguments as-is.
  - Through the pipeline config: Add arguments as environment variables (`env.VAR_NAME`).
2. **Find Environment Variables:**
  - Open a pipeline job in Jenkins, view as **plain text**, and check all available variables.
  - Alternatively - view environment variables: <https://jenkins-git/env-vars.html>.

### 3. Push Changes to Git:

- Create a branch in Skybox named after the Jira ticket (e.g., SKY-272148).
- In `skybox/server/parent/pom.xml`, add a comment to trigger the pipeline.
- Submit a merge request.

### Step 4: Elasticsearch Setup

1. Access Kibana: <https://sb-kibana>. password from secert server - sb-elasticsearch

2. Add the testing index:

- Go to **Management → Dev Tools → Console**.
- Use the following command to create the index:

```
PUT junit_yalmaliah_test
```

- Use the following command to delete the index (needs to be recreated for each run of the pipeline):

```
DELETE junit_yalmaliah_test
```

3. Confirm the document format:

```
GET junit_yalmaliah_test/_search
```

### Step 5: Update Grafana Dashboard

1. Access Grafana: <https://sb-grafana>. password from secert server - Devopsadmin **or** my user (AD)

2. **Duplicate the Dashboard:** Never edit the original.

3. **Add New Variables:**

- Go to the variable bar → **Settings → Variables**.

4. **Update Queries:**

- Edit components → Add `AND x: $x` to the query.

### Step 6: Troubleshooting

- Debug script arguments, by adding to `junit_parser`:

```
print("Parsed options:", opts)
print("Remaining arguments:", args)
```

- Add timestamps to pipeline if needed:

```
script { def currentTime = sh(script: "date '+%H:%M:%S'", returnStdout: true).trim() echo "Timestamp: ${currentTime}" sh label:
'', script: "" python3 pipelineUtils/junit_parser.py -m ${gitlabMergeRequestId} -t ${gitlab_api} \ -p ${WORKSPACE}/skyboxview/
-u ${db_user} -x ${db_pass} -s ${gitlabSourceBranch} \ -d ${gitlabTargetBranch} -o ${gitlabUserEmail} -e ${elastic_pass} -i
${BUILD_ID} \ -r ${gitlabMergeRequestId} -j ${JOB_NAME} -a ${currentTime} "" }
```